

IPFS CONSTELLATION

The Enterprise Control Plane for Decentralized Storage

TECHNICAL WHITEPAPER

Version 1.0

Written on: January 6, 2026

Orchestration • Live Nodes • IPFS CID Tracker

Authors: Nathan Wosnack (Founder & CEO) | MJ Phillips (CTO/CISO)

UbitquityX

A Division of UBITQUITY, INC.

© 2026 UBITQUITY, INC

This page is intentionally left blank.

1. Executive Summary

IPFS Constellation represents a paradigm shift in how enterprises approach decentralized storage infrastructure. As organizations increasingly adopt Web3 technologies and require censorship-resistant, highly available data storage, the limitations of standalone IPFS nodes become critical bottlenecks. Constellation solves this by transforming isolated IPFS nodes into a unified, resilient fleet with military-grade redundancy and automated healing capabilities.

Built on battle-tested [Libp2p](#) and [CRDT \(Conflict-free Replicated Data Types\)](#) consensus protocols, Constellation delivers enterprise-grade data persistence without introducing single points of failure. The system achieves 99.9% data availability through intelligent fleet balancing, geo-aware distribution, and autonomic self-healing mechanisms that operate without human intervention.

Key differentiators include zero-friction adoption through a proxy API identical to standard IPFS, hyperscale ingestion capabilities handling hundreds of pins per second, and granular governance through distinct peer roles. Whether deploying across two nodes or two hundred, Constellation ensures your decentralized storage infrastructure operates with the reliability enterprises demand.

Key Performance Metrics

Data Availability	Single Points of Failure	Geo-Aware Distribution
99.9%	ZERO	100%

2. The Problem: IPFS at Enterprise Scale

2.1 The Promise and Limitations of IPFS

The [InterPlanetary File System \(IPFS\)](#) has emerged as the *de facto* standard for decentralized content storage and distribution. Its content-addressed architecture, where files are identified by their cryptographic hash (CID) rather than location, provides inherent benefits including data integrity verification, deduplication, and censorship resistance. For Web3 applications, NFT metadata storage, and decentralized applications, IPFS offers a compelling alternative to traditional centralized storage.

However, running IPFS at enterprise scale reveals significant operational challenges that the base protocol was not designed to address. While IPFS excels at content distribution and peer-to-peer file sharing, it provides no native mechanisms for ensuring data persistence, managing storage capacity across multiple nodes, or maintaining high availability guarantees that enterprises require.

2.2 Critical Infrastructure Gaps

Single Node Fragility: A standalone IPFS node represents a single point of failure. Hardware failures, network outages, or software crashes can render pinned content completely unavailable. For applications where data availability is mission-critical, this risk is unacceptable.

Manual Replication Burden: Achieving redundancy with standard IPFS requires manually pinning content to multiple nodes and continuously monitoring their status. This operational overhead scales poorly and is prone to human error.

Capacity Management Complexity: As storage requirements grow, administrators must manually track disk utilization across nodes and decide where to route new content. Without intelligent load balancing, nodes become unevenly utilized, leading to premature capacity exhaustion on some nodes while others remain underutilized.

Geographic Distribution Challenges: Global applications benefit from geographically distributed storage to minimize latency. Standard IPFS provides no mechanisms for geo-aware pinning or ensuring content availability in specific regions.

Ingestion Bottlenecks: High-throughput applications generating large volumes of content quickly overwhelm single-node architectures. The lack of queue management and prioritization causes pipeline congestion under heavy load.

2.3 The Enterprise Imperative

Organizations adopting IPFS for production workloads require the same operational guarantees they expect from any enterprise infrastructure: high availability, automatic failover, capacity planning, monitoring, and governance controls. Constellation was designed from the ground up to deliver these capabilities while preserving the decentralized, censorship-resistant properties that make IPFS valuable.

3. Technical Architecture

3.1 Design Philosophy

IPFS Constellation operates as a proactive, self-healing orchestration layer that sits atop existing IPFS daemons. Rather than replacing or modifying IPFS itself, Constellation coordinates multiple IPFS nodes into a unified cluster while each node continues running the standard IPFS implementation. This design ensures compatibility with the broader IPFS ecosystem while adding enterprise-grade coordination capabilities.

The architecture adheres to three core principles:

1. **Leaderless Consensus:** No single node acts as a central coordinator. All nodes participate equally in maintaining cluster state, eliminating single points of failure in the control plane.
2. **Eventual Consistency:** Using CRDTs, the cluster achieves consensus without requiring synchronous coordination, enabling operation across high-latency network links and during network partitions.
3. **Non-Invasive Integration:** Constellation communicates with IPFS nodes through their standard HTTP API, requiring no modifications to IPFS itself and ensuring forward compatibility.

3.2 Component Architecture

Constellation Daemon: Each node runs a Constellation daemon alongside its IPFS daemon. The Constellation daemon manages cluster membership, state synchronization, and local pin operations. Daemons communicate with each other over Libp2p, leveraging the same peer-to-peer networking stack that powers IPFS itself.

Pinset State: The cluster maintains a shared pinset representing all content that should be pinned across the fleet. This state is stored as a CRDT, allowing concurrent modifications from any node without coordination. Conflicts are resolved automatically using commutative, associative, and idempotent merge operations.

Allocation Engine: When content is added to the cluster pinset, the allocation engine determines which nodes should pin the content based on the configured replication factor, current node capacities, and allocation policies (geographic, tagged, etc.). The allocation decision is stored alongside the pin request.

Pin Tracker: Each Constellation daemon runs a pin tracker that monitors the difference between the desired pinset state and the actual pins on the local IPFS node. When discrepancies are detected, the tracker initiates pin or unpin operations as needed.

Proxy API: Constellation exposes an HTTP API endpoint that mirrors the standard IPFS pinning API. Applications can point their existing IPFS tooling at the Constellation proxy, which handles routing requests to the appropriate cluster operations. This enables zero-code-change adoption for existing IPFS applications.

3.3 Consensus and State Synchronization

Constellation employs Conflict-free Replicated Data Types (CRDTs) for maintaining consistent cluster state without requiring a consensus leader. CRDTs are data structures that can be modified concurrently by multiple actors and guarantee that all replicas will converge to the same state, regardless of the order in which updates are received.

The primary data structures include:

- **Observed-Remove Set (OR-Set):** Tracks the global pinset, supporting both add and remove operations with automatic conflict resolution
- **Last-Writer-Wins Register (LWW-Register):** Stores pin metadata and allocation decisions, using logical timestamps to resolve concurrent writes
- **Grow-Only Counter (G-Counter):** Tracks cluster-wide metrics such as total pins and storage utilization

State updates propagate through the Libp2p pubsub system, with each node subscribing to cluster state topics. When a node receives an update, it merges the incoming state with its local state using the CRDT merge operation. Periodically, nodes also perform full state synchronization with peers to recover from any missed updates.

4. Core Capabilities

4.1 Autonomic Data Persistence

Constellation delivers true "set and forget" reliability for pinned content. When a CID is added to the cluster, Constellation asynchronously manages its complete lifecycle:

- **Automatic Retry:** Failed pin operations are automatically retried with exponential backoff. Transient failures (network issues, temporary node unavailability) are handled transparently.
- **Replication Factor Maintenance:** The cluster continuously monitors actual replication levels. If a node becomes unavailable or disk fails, content is automatically re-replicated to maintain the configured redundancy.
- **Health-Based Routing:** Pin operations are routed away from unhealthy nodes, and content is proactively migrated from nodes exhibiting degraded performance.

4.2 Smart Fleet Balancing

The intelligent allocation engine eliminates manual capacity management by dynamically routing storage requests based on multiple factors:

- **Available Capacity:** Nodes with more free space receive proportionally more allocations, preventing premature exhaustion of individual nodes.
- **Geographic Region:** Content can be pinned to specific geographic regions to minimize latency for regional users or comply with data residency requirements.
- **Custom Tags:** Nodes can be tagged with arbitrary labels (e.g., "high-performance", "archival", "customer-A"), and allocation policies can route content to nodes matching specific tags.
- **Rebalancing:** When new nodes join the cluster or capacity changes, the allocation engine can optionally trigger rebalancing to redistribute content more evenly.

4.3 Leaderless Resilience

Traditional distributed systems often rely on leader election to coordinate operations. While effective, this approach introduces a single point of failure: if the leader becomes unavailable, the cluster may become temporarily inoperative until a new leader is elected.

Constellation eliminates this vulnerability through its CRDT-based architecture. Every node maintains a complete copy of cluster state and can accept operations independently. Benefits include:

- **Zero Downtime:** No leader election delays. The cluster continues operating even as individual nodes join, leave, or fail.
- **Partition Tolerance:** During network partitions, each partition continues operating independently. When connectivity is restored, states merge automatically.

- **Horizontal Scalability:** Adding nodes increases cluster capacity without requiring reconfiguration of existing nodes.

4.4 Hyperscale Ingestion

For high-throughput environments generating large volumes of content, Constellation provides a sophisticated ingestion pipeline:

- **Smart Queue:** Incoming pin requests are queued and processed asynchronously. The queue intelligently prioritizes new content over retry attempts, ensuring fresh data flows through even under heavy load.
- **Parallel Processing:** Multiple pin operations execute concurrently, limited by configurable concurrency settings to prevent overwhelming individual IPFS nodes.
- **Backpressure:** When the cluster is overwhelmed, the API returns appropriate status codes allowing clients to implement retry logic, preventing cascade failures.
- **Throughput:** Benchmarks demonstrate sustained ingestion rates of hundreds of pins per second under optimal conditions.

4.5 Granular Governance

xEnterprise deployments require fine-grained access control. Constellation implements a role-based peer system:

Trusted Peers: Full cluster participants that can modify cluster state (add/remove pins, change configuration). Trusted peers participate in state consensus and can accept client API requests.

Follower Peers: Read-only participants that replicate content but cannot modify cluster state. Followers receive state updates from trusted peers and pin content accordingly, but their IPFS nodes serve content to the network without exposing cluster management capabilities.

This separation enables deployment patterns such as:

- Secure core clusters with read-only edge nodes for content distribution
- Multi-tenant architectures where different organizations manage their own trusted peers
- Hybrid deployments spanning internal infrastructure and cloud providers

4.6 Zero-Friction Adoption

Constellation is designed for immediate adoption without requiring application changes. The proxy API implements the standard IPFS pinning service API specification, allowing existing tools and applications to inherit clustering capabilities by simply changing their endpoint configuration.

Compatible tooling includes:

- IPFS CLI and desktop applications
- NFT.Storage, Web3.Storage, and similar pinning services
- Custom dApps using IPFS HTTP client libraries
- CI/CD pipelines with IPFS integration

5. Deployment Architecture

5.1 Minimum Viable Cluster

A production Constellation deployment requires a minimum of three trusted peers to ensure quorum during network partitions. Each peer requires:

- IPFS daemon (Kubo 0.18+ recommended)
- Constellation daemon
- Network connectivity to other cluster peers (Libp2p)
- Sufficient storage for replicated content

5.2 Recommended Production Architecture

For enterprise deployments, we recommend:

- **Geographic Distribution:** Deploy peers across multiple availability zones or regions to protect against regional outages.
- **Load Balancer:** Place an HTTP load balancer in front of Constellation proxy endpoints to distribute client traffic.
- **Monitoring Stack:** Constellation exposes Prometheus metrics for cluster health, pin status, and performance monitoring.
- **Backup Strategy:** While content is replicated across the cluster, maintain off-cluster backups for disaster recovery.

5.3 Live Public Nodes

UbitquityX operates production Constellation clusters accessible for integration and testing:

Node	Endpoint
Primary Gateway	ubitquityx.com
Secondary Gateway	smartescrow.us

6. Enterprise Use Cases

6.1 NFT and Digital Asset Storage

NFT marketplaces and creators require permanent, highly available storage for metadata and media assets. Constellation ensures that once minted, NFT content remains accessible regardless of individual infrastructure failures. The geo-aware distribution feature enables fast content delivery to global audiences.

6.2 Real Estate Document Archival

Property records, deeds, and title documents benefit from immutable, verifiable storage. Constellation's content-addressed architecture provides built-in integrity verification while the orchestration layer ensures long-term availability. UBITQUITY has leveraged this capability to mint over 10 million real estate records since March 2023.

6.3 Decentralized Application Backends

dApps requiring censorship-resistant data layers can use Constellation as a managed storage backend. The API compatibility means existing IPFS-integrated applications work without modification, while gaining enterprise reliability features.

6.4 Media and Content Distribution

Content creators and media companies can use Constellation for decentralized content distribution. The hyperscale ingestion pipeline handles high-volume uploads, while geographic distribution ensures low-latency delivery to end users.

6.5 Compliance and Audit Trails

Regulated industries can leverage IPFS content addressing for tamper-evident record keeping. Constellation adds the availability and governance controls necessary for compliance workloads, including role-based access and geographic pinning for data residency requirements.

7. The AETHER Token Economy

To ensure the economic sustainability and security of the Constellation ecosystem, the AETHER utility token (\$AETHERT) serves as the fuel for the enterprise control plane. Built on the XPR Network, \$AETHERT leverages the unique zero-gas and on-chain identity capabilities of the Antelope protocol to deliver a seamless, high-performance utility layer.

7.1 Token Profile

- Name: AETHER
- Ticker: \$AETHERT
- Chain: XPR Network
- Total Supply: 111,111,111 AETHERT
- Launch Price: \$0.0011 per AETHERT*
- Precision: 8 Decimals

(* *Launching on Alcor Exchange on January 7, 2026 by 11:11 PM ET*)

7.2 Core Utility Pillars

The AETHER token is not merely a payment method; it is the cryptographic binding agent for the decentralized fleet.

- **Identity-Verified Staking (The Trust Bond):** Leveraging XPR's native Verified Identity (KYB) layer, only nodes with verified business status can stake AETHERT to become "Trusted Peers." This stake acts as a security bond; if a node fails to meet the 99.9% uptime SLA, the stake is slashed, ensuring "military-grade" reliability.
- **Resource Provisioning (RAM & CPU):** AETHERT tokens are utilized to provision blockchain RAM for the "Cluster Pinset" state (CRDTs) on-chain. This ensures the control plane remains permanent and decentralized without relying on centralized databases.
- **Micro-Settlement (Pay-Per-Pin):** Leveraging XPR's zero-gas architecture, the network supports streaming payments. Enterprises can pay for storage in real-time (per second) without the friction of transaction fees.

- **Atomic Federation:** As the network scales to multi-cluster federation, AETHERT acts as the automated settlement layer for storage resource exchange between different organizations and future Filecoin cold-storage deals.

7.3 Token Allocation & Distribution

The distribution of AETHERT is designed to incentivize long-term infrastructure stability over short-term speculation, with specific allocations for both node operators and the community.

Category	Allocation	Amount (AETHERT)	Purpose
Node Mining Rewards	25%	27,777,777.75	Incentives for Trusted Peers maintaining 99.9% uptime.
nDAO Treasury Stored on nftitledao account	20%	22,222,222.20	Future growth, development, and emergency fund.
Team & Founder	15%	16,666,666.65	Fair compensation for efforts pertaining to this initiative.
Regular Staking Pool	15%	16,666,666.65	Passive APY for community holders (non-node operators).

<u>Private Sale</u>	15%	16,666,666.65	Initial capital for operations, development, marketing efforts.
Liquidity Farming/Pool on Alcor	10%	11,111,111.1	Rewards for Liquidity Providers on Alcor Exchange.

7.4 Staking Pools

- Node Rewards (25%): Paid exclusively to operators of Trusted Peers who prove valid storage and uptime. This pool emits rewards over a 10-year schedule to guarantee long-term infrastructure.
- Regular Staking (15%): A smart staking pool for the broader community. While these users do not run nodes, their stake can be delegated to Trusted Peers to vote on governance proposals via our partnership at [nDAO](#). Tiers include Flexible (No Lock) and Long-Term (12-Month Lock) with higher APY and "Scout" voting badges.

Staking on BlastPad launches on or before January 9, 2026.

8. Future Development Roadmap

IPFS Constellation continues to evolve with planned enhancements including:

- **Filecoin Integration:** Automated deal-making with Filecoin storage providers for cold storage tier and additional redundancy.
- **Enhanced Monitoring Dashboard:** Web-based UI for cluster visualization, pin management, and operational metrics.
- **Encryption at Rest:** Optional client-side encryption with key management integration for sensitive content.
- **Cost Optimization Engine:** Intelligent tiering between hot (SSD) and cold (HDD/Filecoin) storage based on access patterns.
- **Multi-Cluster Federation:** Ability to federate multiple Constellation clusters for organizational boundaries while maintaining shared content addressing.

9. Conclusion

IPFS Constellation represents a critical infrastructure layer for organizations serious about decentralized storage. By transforming isolated IPFS nodes into coordinated, self-healing fleets, Constellation delivers the reliability and operational characteristics enterprises demand without sacrificing the censorship resistance and decentralization benefits that make IPFS compelling.

Key advantages include:

- 99.9% data availability through automated replication and healing
- Zero single points of failure via leaderless CRDT-based consensus
- Intelligent fleet management with geo-aware and capacity-based allocation
- Hyperscale ingestion for high-throughput applications
- Zero-friction adoption with standard IPFS API compatibility

As Web3 adoption accelerates and organizations require robust decentralized infrastructure, Constellation provides the enterprise-grade foundation for the next generation of content-addressed storage systems.

Contact Information

For deployment inquiries, partnership opportunities, or technical questions:

- **Contact Us:** support.ubitquityx.com (Open a Support Ticket)
- **Constellation Portal:** ubitquityx.com/IPFS_Constellation
- **Documentation:** ubitquityx.com/IPFS_Constellation/docs
- **GitHub:** https://github.com/ubitquity/IPFS-CONSTELLATION_PUBLIC

IPFS Constellation

© 2026 UbitquityX (a division of UBITQUITY, INC.)
ONE BLOCK AT A TIME®